

Using Bayesian Learning to Estimate How Hot an Execution Path is

Karim Ali

{karim}@cs.uwaterloo.ca

David R. Cheriton School of Computer Science
University of Waterloo

Abstract—The abstract goes here.



1 INTRODUCTION

2 MOTIVATION

3 RELATED WORK

4 IMPLEMENTATION

4.1 Path Description Model

4.2 Path Enumeration

4.3 Dataset Generation

5 EXPERIMENTAL RESULTS

6 DISCUSSION AND FUTURE WORK

7 CONCLUSION

REFERENCES

- [1] C. Booger and L. Moonen, “On the Use of Data Flow Analysis in Static Profiling,” in *Source Code Analysis and Manipulation, 2008 Eighth IEEE International Working Conference on*, IEEE, 2008, pp. 79–88.
- [2] S. Heckman and L. Williams, “A Systematic Literature Review of Actionable Alert Identification Techniques for Automated Static Code Analysis.”
- [3] R. Buse and W. Weimer, “The road not taken: Estimating path execution frequency statically,” in *Proceedings of the 2009 IEEE 31st International Conference on Software Engineering*, IEEE Computer Society, 2009, pp. 144–154.
- [4] B. Calder, D. Grunwald, M. Jones, D. Lindsay, J. Martin, M. Mozer, and B. Zorn, “Evidence-based static branch prediction using machine learning,” *ACM Transactions on Programming Languages and Systems (TOPLAS)*, vol. 19, no. 1, pp. 188–222, 1997.
- [5] T. Baba, T. Masuho, T. Yokota, and K. Ootsu, “Design of a two-level hot path detector for path-based loop optimizations,” in *Proceedings of the third conference on IASTED International Conference: Advances in Computer Science and Technology*, ACTA Press, 2007, p. 28.
- [6] T. Ball and J. Larus, “Efficient path profiling,” in *micro*. Published by the IEEE Computer Society, 1996, p. 46.
- [7] E. Duesterwald and V. Bala, “Software profiling for hot path prediction: Less is more,” *ACM SIGARCH Computer Architecture News*, vol. 28, no. 5, pp. 202–211, 2000.
- [8] M. Merten, A. Trick, C. George, J. Gyllenhaal, and W. Hwu, “A hardware-driven profiling scheme for identifying program hot spots to support runtime optimization,” in *Proceedings of the 26th annual international symposium on Computer architecture*, IEEE Computer Society, 1999, p. 147.
- [9] G. Ammons and J. Larus, “Improving data-flow analysis with path profiles,” in *Proceedings of the ACM SIGPLAN 1998 conference on Programming language design and implementation*, ACM, 1998, pp. 72–84.
- [10] S. Graham, P. Kessler, and M. Mckusick, “Gprof: A call graph execution profiler,” *ACM Sigplan Notices*, vol. 17, no. 6, pp. 120–126, 1982.
- [11] W. Weimer and G. Necula, “Exceptional situations and program reliability,” *ACM Transactions on Programming Languages and Systems (TOPLAS)*, vol. 30, no. 2, pp. 1–51, 2008.
- [12] J. Larus, “Whole program paths,” in *Proceedings of the ACM SIGPLAN 1999 conference on Programming language design and implementation*, ACM, 1999, pp. 259–269.
- [13] T. Ball and J. Larus, “Branch prediction for free,” *ACM SIGPLAN Notices*, vol. 28, no. 6, pp. 300–313, 1993.
- [14] P. Mehra and B. Wah, “Synthetic Workload Generation for Load-Balancing Experiments,” *IEEE Parallel & Distributed Technology: Systems & Technology*, vol. 3, no. 3, pp. 4–19, 1995.
- [15] A. Nanda and L. Ni, “Benchmark workload generation and performance characterization of multiprocessors,” in *Proceedings of the 1992 ACM/IEEE conference on Supercomputing*, IEEE Computer Society Press, 1992, pp. 20–29.
- [16] K. Sen, “Concolic testing,” in *Proceedings of the twenty-second IEEE/ACM international conference on Automated software engineering*, ACM, 2007, pp. 571–572.
- [17] K. Sen and G. Agha, “CUTE and jCUTE: Concolic unit testing and explicit path model-checking tools,” in *Computer Aided Verification*, Springer, 2006, pp. 419–423.
- [18] S. Goldsmith, A. Aiken, and D. Wilkerson, “Measuring empirical computational complexity,” in *Proceedings of the the 6th joint meeting of the European software engineering conference and the ACM SIGSOFT symposium on The foundations of software engineering*, ACM, 2007, pp. 395–404.
- [19] R. Vallée-Rai, P. Co, E. Gagnon, L. Hendren, P. Lam, and

- V. Sundaresan, “Soot-a Java bytecode optimization framework,” in *Proceedings of the 1999 conference of the Centre for Advanced Studies on Collaborative research*. IBM Press, 1999, p. 13.
- [20] M. Hall, E. Frank, G. Holmes, B. Pfahringer, P. Reutemann, and I. Witten, “The WEKA data mining software: An update,” *ACM SIGKDD Explorations Newsletter*, vol. 11, no. 1, pp. 10–18, 2009.